

Linux Synchronization Primitives (cont.)

Atomic types

spinlocks = kernel/sched/core.c (see all uses of spinlock)

DEFINE_SPINLOCK (name) is the static way to initialize ^{spinlock_irqsave} for when structure can be modified from an interrupt.

More or Less implemented as:

```
while (test-and-set(x));
```

But with some additions to resolve cases when you grab a lock and then someone preempts you and just sits there spinning and waiting for your lock, wasting time:

```
spinlock(x) {
```

```
    disable_preemption();
```

```
    while (test-and-set(x)) {
```

```
        enable_preemption();
```

```
        disable_preemption();
```

```
    }
    enable_preemption();
}
```

to give a chance to be preempted

in actual Linux interpretation this is probably not here, i.e.

you can't be preempted while holding a spinlock.

To make this an uninterruptible, just disable interrupts at start and re-enable or restore them before returning.

Locks only work if you use them correctly

Monitors are meant to make life easier:

```
monitor {
```

```
    f1() { ... }; f2() { ... }; f3() { ... };
```

```
}
```

By guaranteeing mutual exclusion b/w f1...f3 w/in the monitor.

Example:

```
monitor {
```

```
    float balance;
```

```
    void credit(float a) { balance += a; };
```

```
    void debit(float a) { balance -= a; };
```

```
}
```

only one function can be executing at a time.

Condition Variables - wait for events within the monitor:

```
condition variable {
```

```
    void wait();
```

```
    void signal();
```

```
}
```

(and get out of the monitor)

wait till someone signals

signal the event

NOTE: This allows threads to do things that depend on one another

Suppose t_0 is in wait() and t_1 does signal(), now t_1 is in the monitor. But t_0 wants to get in - which one do we run? Either way can be implemented. Hoare says run t_0 immediately & suspend t_1 . Hansen says to keep running t_1 . There is no right or wrong answer.

► Monitors don't exist as a construct in the Linux kernel.

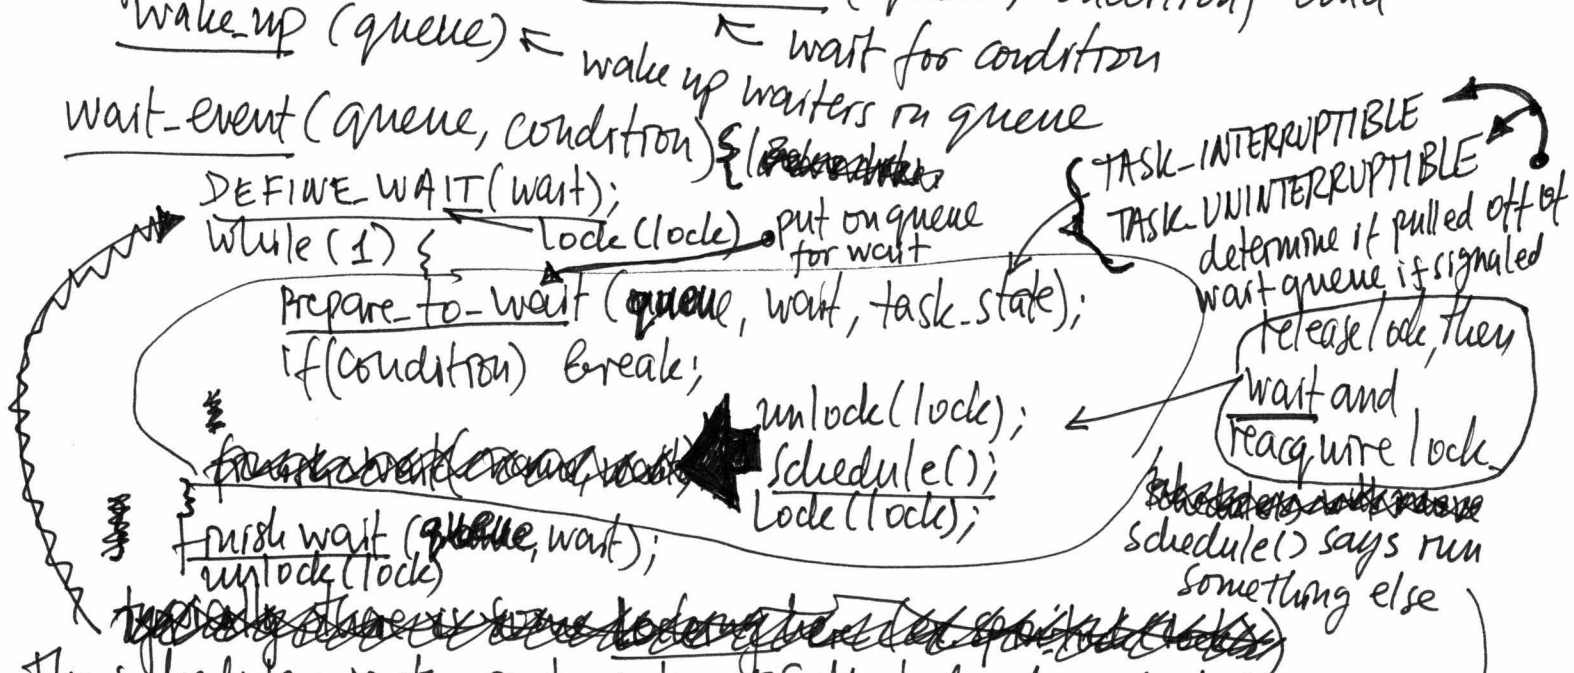
But the idea of a critical section with signaling (condition variables) goes beyond monitors, for example:

- Critical section implemented with a lock
- Within it wait for a condition/event, giving up the lock at the same time
- When done waiting, grab the lock again & continue

Basically, the core idea of using a condition variable w/in a critical section.

► Waitqueues are Linux's mechanism for going from running to wait for some event and becoming nonrunnable

► High-level constructs are wait-event (queue, condition) and wake-up (queue)



The scheduler is the part of the OS that decides what to run on the CPU (typically, triggered by a timer interrupt).

/kernel/sched/core.c → --schedule(void)

You keep ~~waiting~~ ^{sleeping} till someone calls wake-up(), which will find jobs on the wait queue and change their state to **TASK-RUNNING** (meaning task is runnable, Linux does not distinguish b/n running & runnable). Then the scheduler picks it up from there. Processes that were waiting pick up from schedule() onwards.

Linux has a Runqueue (a List of runnable/running tasks). Tasks that are waiting are on a waitqueue, but when are signaled or woken up, they get moved to the runqueue w/ status TASK_RUNNABLE. The scheduler only looks at the runqueue.

- wakeup_all() wakes up all waiters, wakeup() wakes up one
- wait_event() is uninterruptible, wait_event_interruptible() is interruptible

► For more complex wait checks (i.e. when a simple condition is not sufficient), we can just code sth. similar to the wait_event() implementation. or search for examples how prepare-to-wait() gets used. ↓

pretty much what we have on prev. ^{Page.} include/linux/wait.h

- Use add_wait_queue() after DEFINE_WAIT(wait) to add the wait together w/ other waits in the same queue (optional?)
- usually wait in a loop to re/define condition and sleep again, if necc.
- prepare-to-wait() is to make yourself not running (in a blocked state) and add yourself to the queue.
- Reference Counting - check LKD book for description
Similar to RCU in that you keep track of users for sth. you need to release.
- Kfifo - kernel FIFO Buffer
- IDR - integer id management mechanism

● Scheduler - A thread runs, at some point, because of an interrupt (HW or syscall) you get into the kernel. Till you get to the kernel, this thread will always be running. Once you are in the kernel code, it will check if it needs to run something else. That happens because somewhere in the OS, scheduler() gets called:

- either directly, like in wait_event() to release processor
- or indirectly by setting a schedule flag in interrupt code (for ex.) since we don't want to call schedule (a complex call) there, and then (after interrupt is done) checking it and calling scheduler() if necessary.