

HW6 = no demos!

## ● Storage / Disk

Primary abstraction is Files, supported by a File System - key issues are performance & reliability.

► Allocation is also an important issue:

- Linked approach - Both performance & reliability are poor
- FAT - replication gives good reliability, caching gives good performance
- Indexed is what Linux uses; AKA index node or inode

Index (per file, fixed size)



- Each file has an index table w/ each index entry pointing to a block entry of the file. (size of index is fixed)
- Index is crucial for reliability, but it is small (can be cached) and per file (can't store all files)
- Index is stored on disk @ special location.

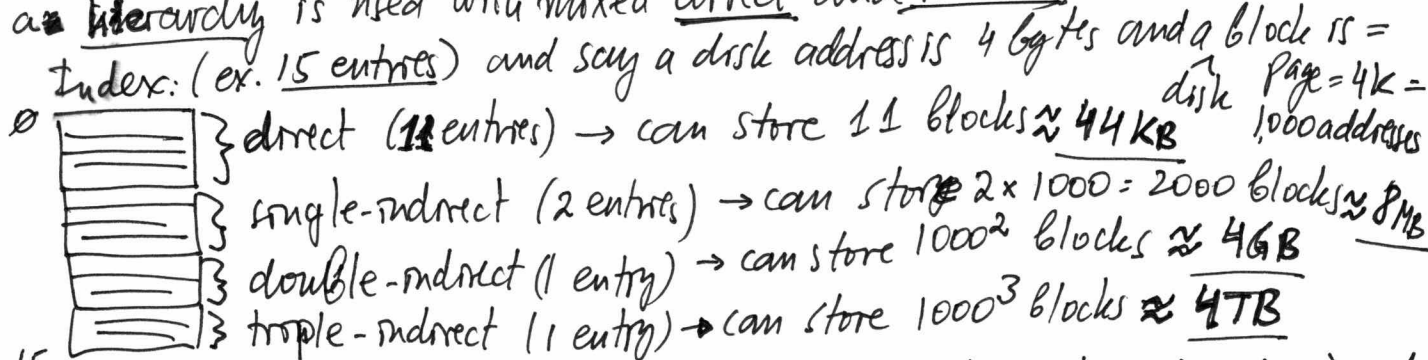
► Block = fixed size unit of storage on the disk (often = frame) ← typically blocks don't contain multiple files.

► Real file systems are not strictly trees, since links b/n files are cycles.

► symlink = removing the link leaves the file; hard link = removing link removes file.

► inode size is fixed (typically); block size is fixed too! typically, an inode has around 15 entries; block size that is large will result in wasted space (most files will be much smaller & can't share blocks)

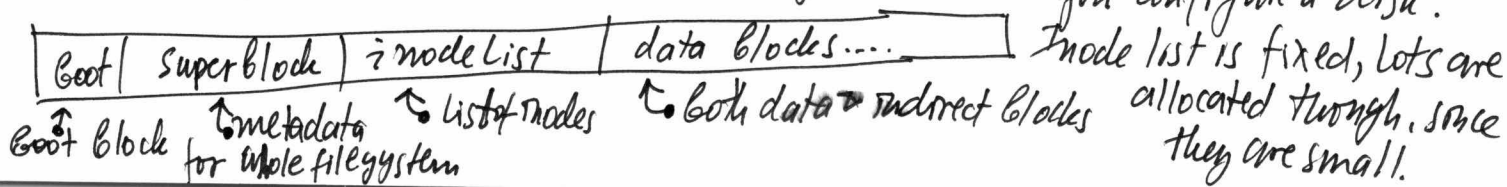
► So a hierarchy is used with mixed direct and indirect index entries:




**NOTE:** The inode is NOT a disk block, it's a separate (much smaller structure) and inodes are pre-allocated (in certain disk blocks). The indirect blocks are special blocks that are interpreted as arrays of (indirect or direct) addresses.

**Ex 1:** UNIX System V (SFS): inode = 64 bytes  
block = 512 bytes


everything gets setup when you configure a disk.



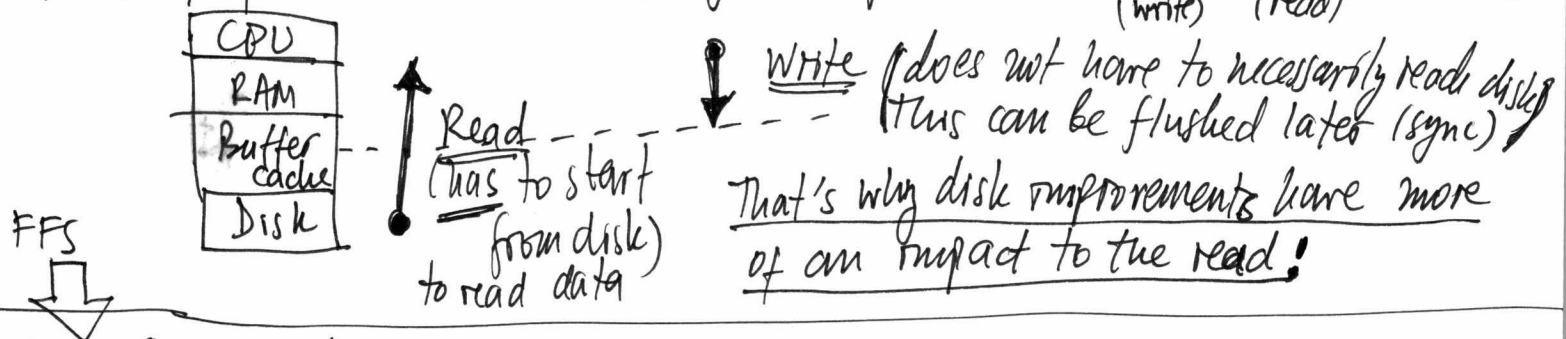
EX1 (cont.)

But, this implementation of the index scheme is not used widely anymore (the index scheme is still used, but w/ different implementation).  
 Performance - inodes & data blocks are separated, need to go back and forth  
 Reliability - FS metadata and inodes are contiguous and single copy, if problem then the whole FS or most of the files will be lost.

(EX2) FFS (Fast File System) in BSD (Berkeley Standard Distribution) (in metadata)



- ① Superblock is replicated = better reliability
  - ② Locality is increased, since inodes & data blocks don't span the whole disk
- both performance and reliability are improved (about 5 to 8 times better!)  
 (write) (read)



• This became the Linux ext 2. → ext 3 → ext 4

(Ex.3) Journaling (logging) added (w/o journaling, file system checks done by fsck on ext2 take a long time.)

- In addition to storing the data normally, you add a journal / log
- The log / journal is linearly ordered by time = never go back, just keep adding changes
- You can recover the entire file system from the log and no need to look at the entire file system, just the log.